

Web Programming

Databases 3: Constraints, Normalization, Indexing & Transactions

Part III

Constraint

Learning Objectives

- Understand and apply SQL constraints.
- Learn about foreign keys and their role in data integrity.
- Gain an understanding of database normalization.
- Learn to connect multiple tables using joins.
- Briefly discuss indexing and transactions.

What are Constraints?

Constraints enforce rules to maintain data integrity.

Types of Constraints:

- **NOT NULL**: Prevents null values in a column.
- **UNIQUE**: Ensures unique values.
- **DEFAULT**: Assigns default values.
- **CHECK**: Validates data based on a condition.
- **PRIMARY KEY**: Uniquely identifies a record.
- **FOREIGN KEY**: Links tables together.

What are Constraints?

Constraints enforce rules to maintain data integrity.

Types of Constraints:

- **NOT NULL**: Prevents null values in a column.

```
CREATE TABLE students (  
  id INTEGER PRIMARY KEY,  
  name TEXT NOT NULL  
);
```

What are Constraints?

Constraints enforce rules to maintain data integrity.

Types of Constraints:

- **UNIQUE**: Ensures unique values.

```
CREATE TABLE users (  
  id INTEGER PRIMARY KEY,  
  email TEXT UNIQUE  
);
```

What are Constraints?

Constraints enforce rules to maintain data integrity.

Types of Constraints:

- **DEFAULT:** Assigns default values.

```
CREATE TABLE orders (  
  id INTEGER PRIMARY KEY,  
  status TEXT DEFAULT 'Pending'  
);
```

What are Constraints?

Constraints enforce rules to maintain data integrity.

Types of Constraints:

- **CHECK**: Validates data based on a condition.

```
CREATE TABLE employees (  
  id INTEGER PRIMARY KEY,  
  age INTEGER CHECK(age >= 18)  
);
```


What are Constraints?

Constraints enforce rules to maintain data integrity.

Types of Constraints:

- **PRIMARY KEY:** Uniquely identifies a record.

```
CREATE TABLE departments (  
  id INTEGER PRIMARY KEY,  
  name TEXT NOT NULL  
);
```

What are Constraints?

Constraints enforce rules to maintain data integrity.

Types of Constraints:

- **FOREIGN KEY:** Links tables together.

```
CREATE TABLE orders (  
  id INTEGER PRIMARY KEY,  
  customer_id INTEGER,  
  FOREIGN KEY (customer_id ) REFERENCES costumers(id)  
);
```

Column Constraints in SQLite

SQLite does not strictly enforce data types but has type affinity.

Example:

```
CREATE TABLE users (  
  id INTEGER PRIMARY KEY,  
  username TEXT UNIQUE NOT NULL,  
  age INTEGER CHECK(age>=18)  
);
```

Exercise: Insert invalid data and analyze errors.

```
INSERT INTO users (id, username, age) VALUES (1, 'JohnDoe', 17); -- Should fail due to CHECK constraint
```

Foreign Keys and Referential Integrity

SQLite does not enforce foreign keys by default.

To enable:

```
INSERT INTO users (id, username, age) VALUES (1, 'JohnDoe', 17); -- Should fail due to CHECK constraint
```

Example of a Foreign key

```
CREATE TABLE orders (  
  id INTEGER PRIMARY KEY,  
  customer_id INTEGER,  
  FOREIGN KEY (customer_id) REFERENCES customers(id)  
);
```

Cascading Actions

What happens if we UPDATE or DELETE an entry from one of the linked tables.

To enable:

```
ON DELETE CASCADE -- (Delete related records automatically)
```

```
ON UPDATE SET NULL
```

Example

```
CREATE TABLE enrollments (  
  student_id INTEGER,  
  course_id INTEGER,  
  PRIMARY KEY (student_id, course_id),  
  FOREIGN KEY(student_id) REFERENCES students(id) ON DELETE CASCADE  
);
```

Part III

Normalization

Database Normalization

Why Normalize?

- Reduces redundancy and improves consistency.
- Improves query performance and maintains data integrity.

Not-Normalized Database Table

STUDENT			
<u>StudentID</u>	StudentName	MajorName	NoOfCreditHours
111	Kirsten	Accounting	152
222	Eve	IS	138
333	Zoe	IS	138
444	Ben	Accounting	152

Normalized Database Tables

MAJOR	
<u>MajorName</u>	NoOfCreditHours
Accounting	152
IS	138

STUDENT		
<u>StudentID</u>	StudentName	MajorName
111	Kirsten	Accounting
222	Eve	IS
333	Zoe	IS
444	Ben	Accounting

First Normal Form (1NF)

Each column contains atomic values (no lists or arrays).

- **Example**

```
CREATE TABLE employees (  
  id INTEGER PRIMARY KEY,  
  name TEXT,  
  skill TEXT -- Bad Design (should be separate tables)  
);
```

- **Better**

```
CREATE TABLE employee_skills (  
  employee_id INTEGER,  
  skill TEXT,  
  FOREIGN KEY (employee_id) REFERENCES employees(id)  
);
```


Second Normal Form (2NF)

Removes partial dependencies.

- **Example:** Split employee table into employee and department tables.

```
CREATE TABLE employees (  
  id INTEGER PRIMARY KEY,  
  name TEXT,  
  department_id INTEGER,  
  FOREIGN KEY (department_id) REFERENCES departments(id)  
);
```

Third Normal Form (3NF)

No transitive dependencies (remove indirect relationships).

- **Example:** Separate city and country from an address table.

```
CREATE TABLE addresses (  
  id INTEGER PRIMARY KEY,  
  city TEXT,  
  country TEXT  
);
```

- **Exercise:** Given a denormalized table, normalize it step by step.

Exercise #1

Given a denormalized table, normalize it step by step.

1. Denormalized Table

```
CREATE TABLE employee_details (  
  id INTEGER PRIMARY KEY,  
  name TEXT,  
  department TEXT,  
  manager TEXT,  
  skills TEXT  
);
```

2. Step 1 (1NF - Remove Multi-Valued Columns)

```
CREATE TABLE employees (  
  id INTEGER PRIMARY KEY,  
  name TEXT,  
  department TEXT,  
  manager TEXT  
);
```

```
CREATE TABLE employee_skills (  
  employee_id INTEGER,  
  skill TEXT,  
  FOREIGN KEY (employee_id) REFERENCES employees(id)  
);
```

Exercise #1

Given a denormalized table, normalize it step by step.

3. Step 2 (2NF - Remove Partial Dependencies):

```
CREATE TABLE departments (  
  id INTEGER PRIMARY KEY,  
  name TEXT,  
  manager TEXT  
);  
ALTER TABLE employees ADD COLUMN department_id INTEGER;  
UPDATE employees SET department_id = (SELECT id FROM departments WHERE name = employees.department);  
ALTER TABLE employees DROP COLUMN department;
```

4. Step 3 (3NF - Remove Transitive Dependencies):

```
CREATE TABLE managers (  
  id INTEGER PRIMARY KEY,  
  name TEXT  
);  
ALTER TABLE departments ADD COLUMN manager_id INTEGER;  
UPDATE departments SET manager_id = (SELECT id FROM managers WHERE name = departments.manager);  
ALTER TABLE departments DROP COLUMN manager;
```

Exercise #1

- Insert sample data into denormalized table and normalize it step by step.

```
INSERT INTO departments (id, name) VALUES (1, 'Engineering'), (2, 'HR');  
INSERT INTO employees (id, name, department_id) VALUES (1, 'Alice', 1), (2, 'Bob', 2);  
INSERT INTO managers (id, name) VALUES (1, 'Charlie'), (2, 'Dana');  
UPDATE departments SET manager_id = 1 WHERE id = 1;  
UPDATE departments SET manager_id = 2 WHERE id = 2;  
INSERT INTO employee_skills (employee_id, skill) VALUES (1, 'Python'), (1, 'SQL'), (2, 'HR Management');
```

- Write queries using INNER and LEFT JOIN to retrieve meaningful insights.
- Compare query performance before and after indexing.
- Simulate a transaction with rollback and commit.

Part III

Indexing in SQL

Indexing

What is Indexing?

- Indexing improves query performance by allowing the database to find rows faster.
- Works like an index in a book: instead of scanning the entire table, the database uses the index to jump directly to the data.

Types of Indexes

- **Primary Index:** Automatically created for primary keys.
- **Unique Index:** Enforces unique values.

```
CREATE UNIQUE INDEX idx_users_email ON users(email);
```

- **Composite Index:** Index on multiple columns.

```
CREATE INDEX idx_employee_dept ON employees(department_id, name);
```

- **Full-Text Index:** Used for searching text fields (not supported in SQLite).

Trade-Offs of Indexing

Pros:

1. Speeds up searches and queries.
2. Enhances efficiency for large datasets.

Cons:

1. Slows down INSERT, UPDATE, and DELETE operations.
2. Takes up additional storage space.

Part III

Additional SQL Topics

Auto-Increment in SQL

Used to automatically generate unique values for a primary key.

Example:

```
CREATE TABLE users (  
  id INTEGER PRIMARY KEY AUTOINCREMENT,  
  name TEXT NOT NULL  
);
```

When inserting a row, SQLite automatically assigns the next available ID.

Handling Dates in SQL

SQLite does not have a dedicated DATE type but supports storing dates as:

- *TEXT* (ISO 8601 format YYYY-MM-DD HH:MM:SS)
- *INTEGER* (Unix timestamp)
- *REAL* (Julian day number)
- **Example**

```
CREATE TABLE events (  
  id INTEGER PRIMARY KEY,  
  event_name TEXT,  
  event_date TEXT DEFAULT CURRENT_TIMESTAMP  
);
```

- Extracting Date Parts

```
SELECT strftime('%Y', event_date) AS year FROM events;
```

SQL Views

Views are virtual tables that simplify complex queries.

- **Example**

```
CREATE VIEW employee_details AS  
SELECT employees.name, departments.name AS department  
FROM employees  
JOIN departments ON employees.department_id = departments.id;
```

- **Querying a View**

```
SELECT * FROM employee_details;
```

SQL Injection and Security

What is SQL Injection?

A technique where malicious SQL statements are inserted into an input field.

- **Example** of a vulnerable query:

```
SELECT * FROM users WHERE username = 'admin' AND password = ' ' OR '1' = '1';
```

Preventing SQL Injection

- Use parameterized queries:

```
SELECT * FROM users WHERE username = ? AND password = ?;
```

Hosting a Database

SQLite is serverless, but other databases like PostgreSQL and MySQL require a database server. Common hosting solutions:

- Local Development: SQLite, MySQL, PostgreSQL.
- Cloud-Based Solutions: AWS RDS, Google Cloud SQL, Azure SQL Database.